

Vijay Saraswat - X10
How do you write code to run on:
50,000 cores
↑
very large number cores

Resilient X10

—
K

~2000 done work

Scala use of concurrency

High performance computing at scale

Question - what if one node goes down

Get 1 construct

↳ automatically handle node failures

- understand failure of nodes
- programmer aware & write compensating behavior

write our own data structures to deal w/ the

APGAS programming model

asynchronous global address space
partitioned

realized natively in X10

extension of APGAS to resilience

Resilient APGAS

alternatives:

(NPI open NP)

operational semantics

resilience is a "slippery" concept

⇒ formal semantics

abstracted version of language

happens before - invariance principle

HBI

X10 - imperative language
problems - mutating state,
arrays

arises issues, memory model, etc.

Managed X10 (on JVM)

better than Spark

Spark - RDD data structure
resilience built on this

Alt research massive asynchrony (global load balancing)
irregular computation on large clusters
richer in X10 than in spark

X10 language

Scala-like syntax
own language (Scala too young)
if redone, would be Scala-based

- richer type system than Java

- constrained type system

- variables before type syntax

(no supports for Scala implicits)

focus on scale program

many nodes
each w/ many threads

parallelism & distribution

focus: productivity

ex - 1000 nodes

- imperative

What kind of coding model give to prog

any prog in any thread can access global data structure

- local
- remote - transferred to other node

4-5 diff orders of magnitude
latency

(DSM - no guarantee about performance)

PGAS - similar

but we added "twist"

↳ Heap is not flat

introduce:

Place

↳ "OS process" (intuition)
memory (virtual
threading
run programs

APGAC

surface places in program model.

worry about places, (data)

- where that state lives?

asynchronous is critical
(change from PGAS)

Single Proc

SPMD Single Program Multiple Data

ex. 2D array

ex. weather prediction
ex. heat transfer across wall of room

large array broken into cells

cell - value point in time

run w/ many points in time

end w/ gradient

write code once

distribute to places

send signal to all places

run till hit Barrier

shuffling of data

ex. wall broken into 100×100 grid

physical boundary adjacent
may be far away node

periodically update boundary condition variable

- local compute
- boundary
- shuffle
- repeat

want SPMD runs in (APGAS) model

want 100,000 nodes
problem to impl efficiently

APGAS

- 1) places
 - 2) asynchronous
- } programmer deals with

designed }
Available } for real work

native or managed ?
(JVM)

new machine GC, threading

(HPC)

(Enterprise)

IBM
(Both!)

(most features work on both)
open-source from beginning

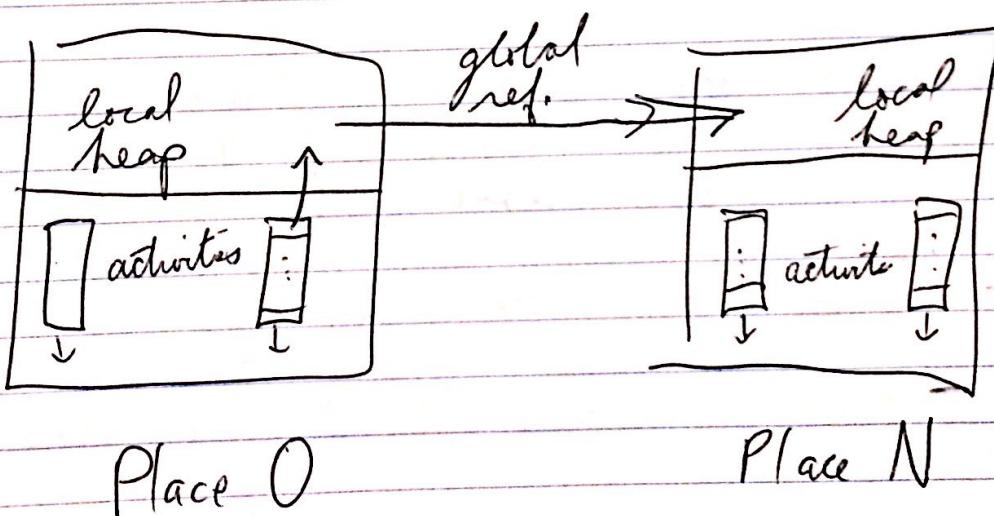
Eclipse - constraints usage
e.g. incremental type-checking

NPI - share nothing abstraction
(no asynchrony)

Open NP - exec large tasks in threads

X - CUDA-like accelerator

Our goal - 1 prog. model



~2004 Task parallelism
· async S
· finish S

updates happen in an order

finish (dual to async)

S - tree of async comp.

run till all are terminated

async - not return "thread" obj
design choice

~~data structures~~

at runtime could be many
or few threads

Place-shifting

at(p) S

- ↳ -stop curr thread
- move state to ' p '
- then exec 'S'

no restriction on nesting

synchronous - waits for S to complete (?)

at(p) e ~ expression-version
returns value

nested structures
code

- refer to outer scope in same way
- leave to runtime to set handle on all references in scope
(read-only)
- only can mutate in its place

Global Ref [T]

new Obj, get local ref
pass to global ref
Generics not impl by erasure (!)
 \Rightarrow I runtime cost

can ask
is pointing to local obj?
if yes, can get it back

Place Local Handle [T]

notion: Global handle freely shared
mutable state, 1 place
PlaceLocal - local obj created in each place
collection

collection of objects

can only get the local obj for a place

Concurrency

(least successful design)

atomic S

exec as if in a single step

no R/W混杂 with this S

S could not have:

async
finish
at

when(c) S

conditional version

(Hoare - cond critie region)

only enter / exec S when c

but both atomic / one step

had to fall back on latches, locks
for performance

(ownership types)

remote val

v = at (p) evalHere (arg1, arg2);

SPMD

finish for (p in Place.places()) {

} at (p) async runEverywhere();

! not
resilient

(no barrier here)

introduce clocks for this

lexically scoped barrier

atomic remote update

at (ref) async atomic ref() += v;

Data Exg.

val _l = l();

finish at (r) async {

val _r = r();

r() = -l;

at (i) async l() = -r;

}

not atomic
w/ resp to
other update

worked for 50,000 places

sequential language

Java w/ Scala - syntax
model

not reusing Java compiler
(need C++ toolchain)
own compiler

global GC
(can have cycles across places)

(global ref counting GC)

Java tricky obj creation
allows "proto" obj to be leaked
↳ hard-hat init

lang rules for in constructor
much stricter

doesn't leak proto-objects

value types - structs
(no mutable state,
just copy freely)

(distinguish copy freely vs not)
struct of

Java array sit off overhead problem

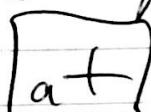
starts → "built-in" types
created by program

ex. complex numbers

kept exceptions → resilience

powerful

resilient x10 - place crashes

only got to place non-zero
by an 

[at] throws exception

thrown async'ly

already built for resilience

Rail vs Array



Java-style

multi-dim
array

distrib arrays